

Bangladesh Data Analysis

Gustavo Cepparo and Milica Cudina

As before, first we import the data.

```
bangladesh=read.csv("bangladesh-data.csv", header=TRUE)
names(bangladesh)

## [1] "Arsenic"  "Chlorine"  "Cobalt"
#accessing single columns
#bangladesh$Arsenic
attach(bangladesh)
Arsenic

##   [1] 2400     6   904    321   1280   151   141   1050   511   688    81     8    37     6    22
##  [16]    43    39    92   253    200   255   1150   1180     9   107     6   149     6    46    13
##  [31]     6   150     6   189    364    42    390     6   270   248   139     6    82    82   256
##  [46]   165     6   180     86     6    38   262    404     8    85    98     6    22     6     6
##  [61]     6    15   103    86     6    46    62    43     6     6    55     6   107    65   276
##  [76]   114     6     6    65   142   194     6    54   702     6   986   153    84    16
##  [91] 1460   306    49    36   106     6    41    84   278     41
##  [ reached 'max' / getOption("max.print") -- omitted 171 entries ]

mean(Arsenic)

## [1] 125.3199

var(Arsenic)

## [1] 88789.39

#hist(Arsenic)
#what are the dimensions of `bangladesh`?
dim(bangladesh)

## [1] 271     3
#see if there are missing data
bangladesh=na.omit(bangladesh)
#what are the "new" dimensions of `bangladesh`?
dim(bangladesh)

## [1] 268     3
attach(bangladesh)

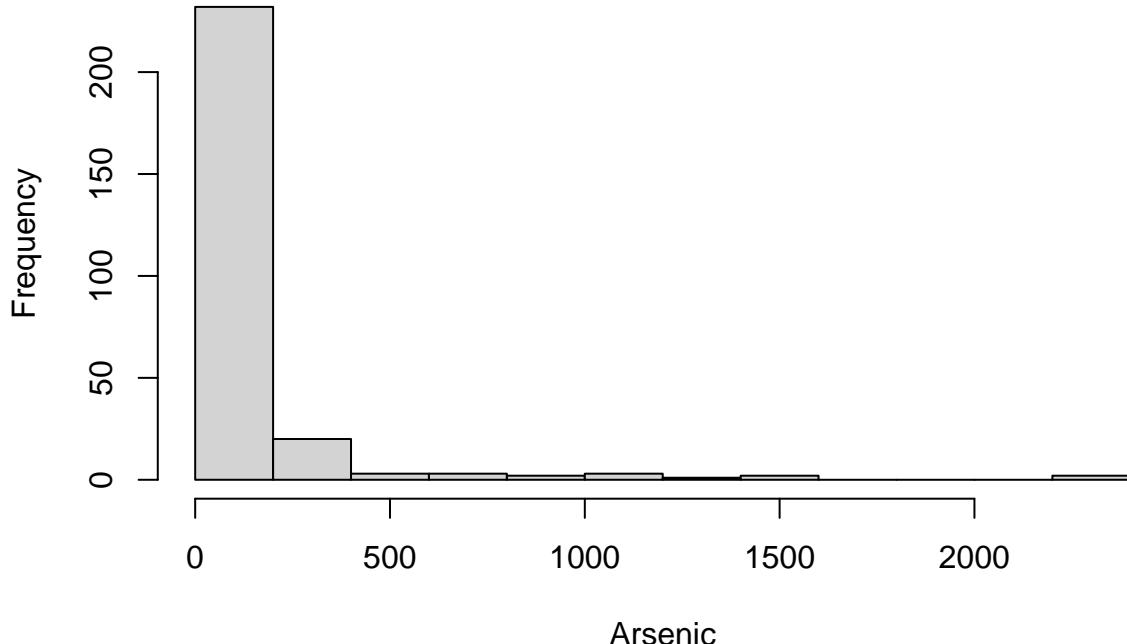
## The following objects are masked from bangladesh (pos = 3):
##
##      Arsenic, Chlorine, Cobalt
n=length(Arsenic)
n
```

```
## [1] 268
```

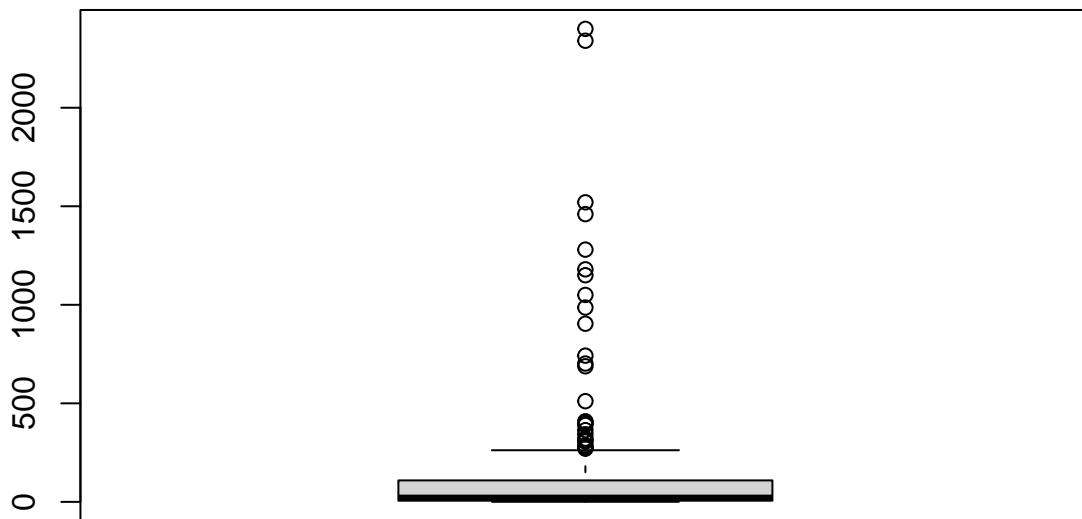
Again, we undertake a rudimentary exploratory data analysis.

```
hist(Arsenic)
```

Histogram of Arsenic



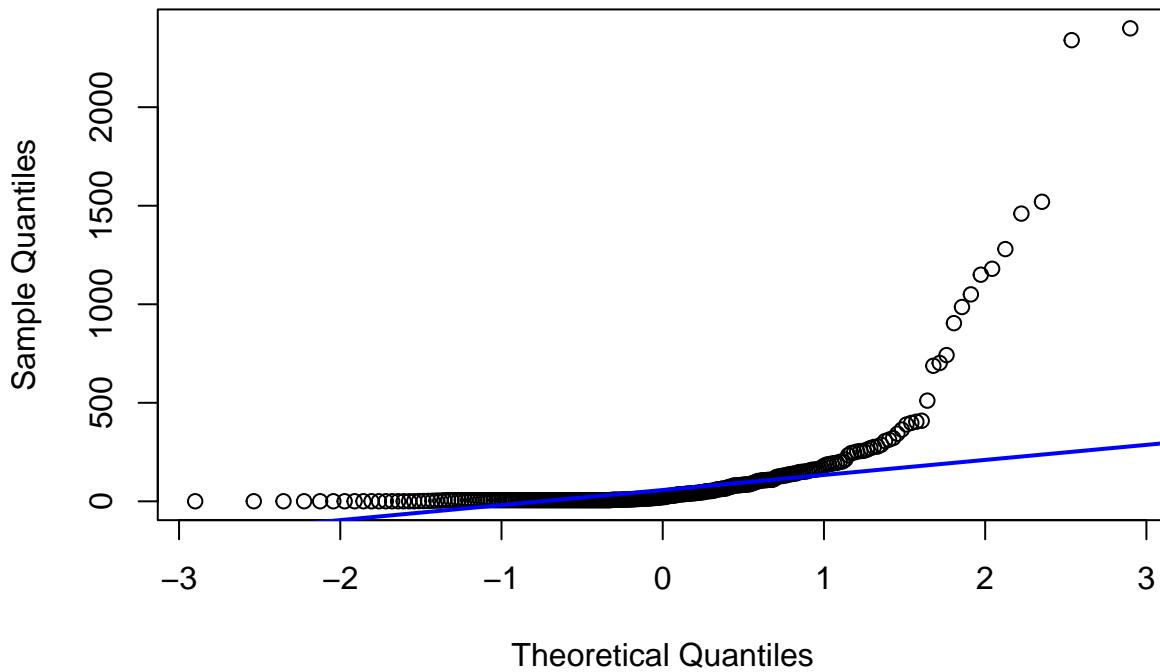
```
boxplot(Arsenic)
```



```
qqnorm(Arsenic)
```

```
qqline(Arsenic, col="blue", lwd=2)
```

Normal Q-Q Plot



What does the test of normality tell us?

```
shapiro.test(Arsenic)
```

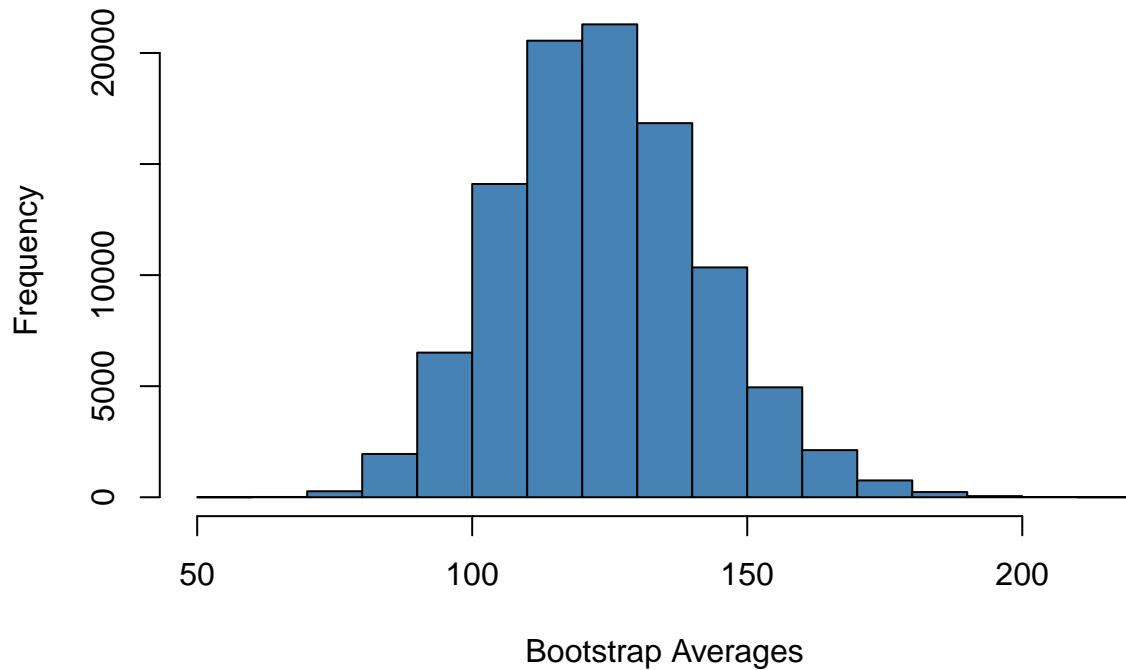
```
##  
## Shapiro-Wilk normality test  
##  
## data: Arsenic  
## W = 0.42284, p-value < 2.2e-16
```

Even though we have ample evidence against the normality of the data, due to the large sample size, we could use the classical approach to the confidence interval.

Now, what about bootstrap?

```
n.boot=10^5  
arsenic.mean=replicate(n.boot, mean(sample(Arsenic, n, replace=TRUE)))  
hist(arsenic.mean,  
     main="Bootstrap Distribution of Averages",  
     xlab="Bootstrap Averages",  
     col="steelblue")
```

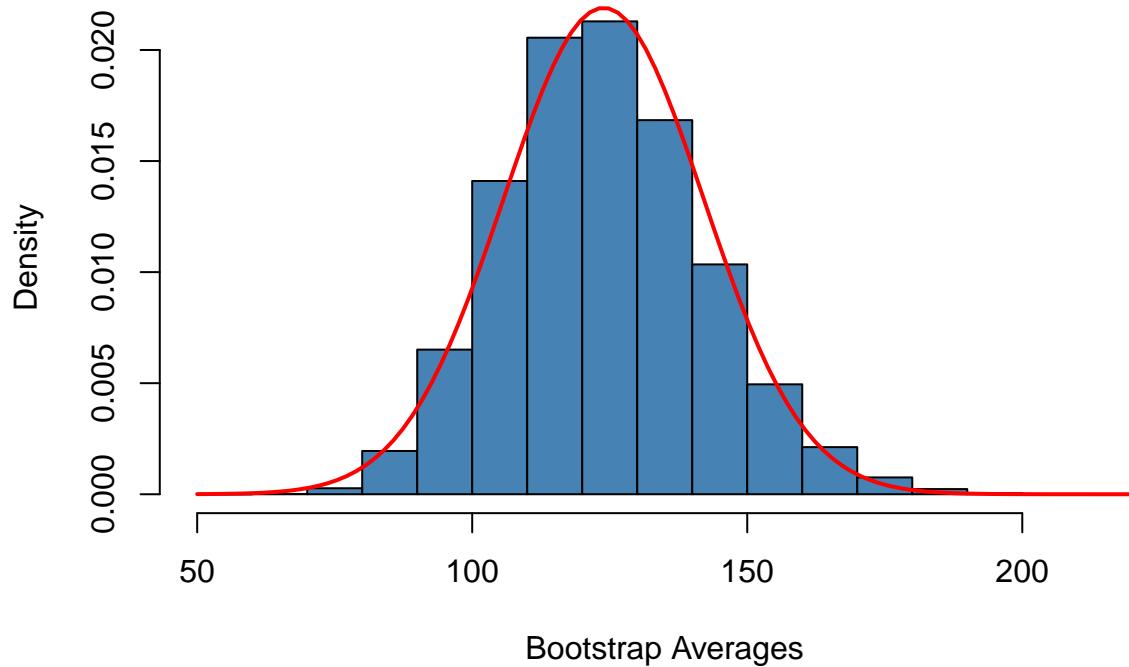
Bootstrap Distribution of Averages



Superimposing the normal bell curve.

```
hist(arsenic.mean,
  main="Bootstrap Distribution of Averages",
  xlab="Bootstrap Averages",
  col="steelblue",
  prob=TRUE)
curve(dnorm(x, mean=mean(arsenic.mean), sd=sd(arsenic.mean)), col="red", lwd=2, add=TRUE)
```

Bootstrap Distribution of Averages



We could now construct a 2SE bootstrap confidence interval.

```
#bootstrap mean
mu.boot=mean(arsenic.mean)
mu.boot

## [1] 123.8937

#bootstrap SE
se.boot=sd(arsenic.mean)

#lower bound
l.bd=mu.boot-2*se.boot
#upper bound
u.bd=mu.boot+2*se.boot

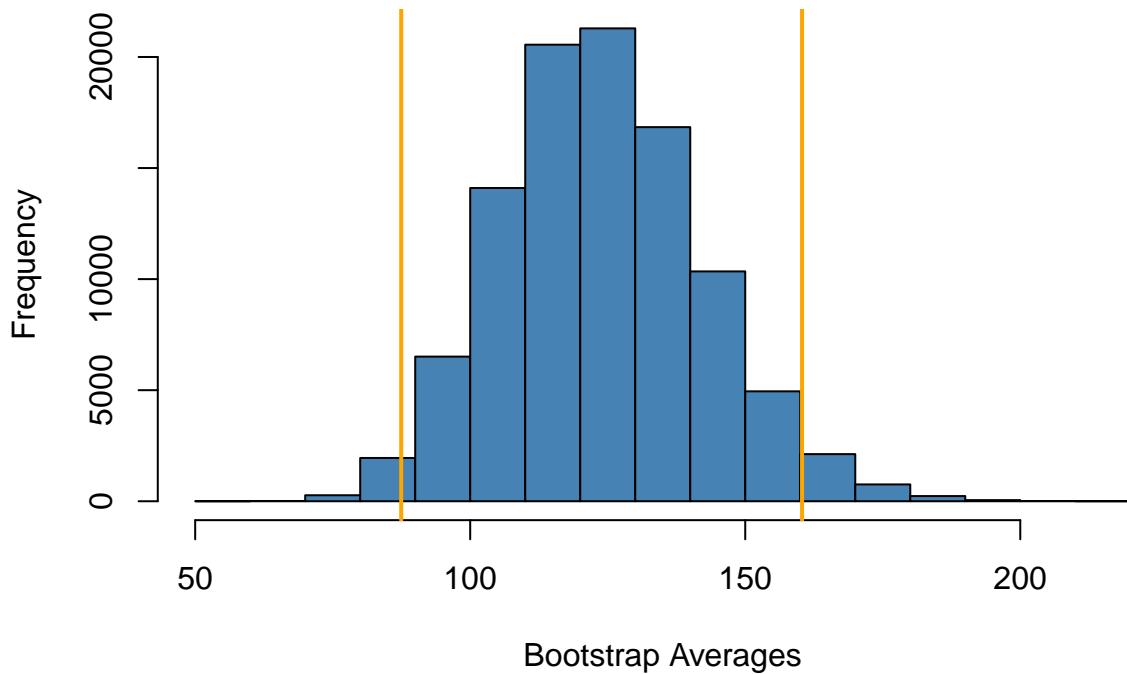
print(c(l.bd, u.bd))

## [1] 87.45872 160.32870
```

It might be interesting to superimpose it on the histogram.

```
hist(arsenic.mean,
  main="Bootstrap Distribution of Averages",
  xlab="Bootstrap Averages",
  col="steelblue")
abline(v=l.bd, col="orange", lwd=2)
abline(v=u.bd, col="orange", lwd=2)
```

Bootstrap Distribution of Averages



We can also construct a bootstrap 95%-percentile confidence interval.

```
median(arsenic.mean)  
  
## [1] 123.02  
bds=quantile(arsenic.mean, c(0.025, 0.975))  
bds  
  
##      2.5%    97.5%  
##  90.75745 162.36031
```

An analogous plot.

```
hist(arsenic.mean,  
     main="Bootstrap Distribution of Averages",  
     xlab="Bootstrap Averages",  
     col="steelblue")  
abline(v=bds, col="orange", lwd=2)
```

Bootstrap Distribution of Averages

