

Logistic regression in medical-school admissions

Gustavo Cepparo and Milica Cudina

The data set we are going to use exists in the `Stat2Data` package. You only need to install it once. Here are the command lines for reference.

```
#install.packages("Stat2Data")
library(Stat2Data)
data(MedGPA)
```

Next, we look at the included data set:

```
data<-MedGPA
#View(data)
data$Acceptance <- as.factor(data$Acceptance)
attach(data)
dim(data)

## [1] 55 11
```

Now, let's see how good the GPA alone is at predicting Acceptance.

```
glm.fits=glm(Acceptance~GPA,data=data,family=binomial)
summary(glm.fits)

##
## Call:
## glm(formula = Acceptance ~ GPA, family = binomial, data = data)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -19.207     5.629  -3.412 0.000644 ***
## GPA          5.454     1.579   3.454 0.000553 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 75.791  on 54  degrees of freedom
## Residual deviance: 56.839  on 53  degrees of freedom
## AIC: 60.839
##
## Number of Fisher Scoring iterations: 4
```

How do we predict with this model?


```
#we'll look at an imaginary student with GPA equal to 3.86
violet=data.frame(GPA=3.86)
violet.prob=predict(glm.fits,violet,type="response")
violet.prob
```

```
##      1
## 0.8637248
```

What about the probability of being accepted as a function of the GPA?

```
mesh=data.frame(GPA=seq(0,4,0.05))
acc.prob=predict(glm.fits, mesh, type="response")
plot(seq(0,4,0.05), acc.prob,
     col="blue", pch=20,
     main="Predicted probability of acceptance",
     xlab="GPA", ylab="Probability")
abline(h=0.5, col="red", lwd=2)
```

Predicted probability of acceptance



Now, what about the confusion matrix?

```
glm.probs=predict(glm.fits, type="response")
glm.pred=rep("yes",55)
glm.pred[glm.probs<.5]="no"
#Confusion Matrix Below
tab=table(glm.pred,Acceptance)
tab

##           Acceptance
## glm.pred  0   1
##       no   16  6
##       yes   9 24
good=(tab[1,1]+tab[2,2])/sum(tab)
good

## [1] 0.7272727
```

Let's "improve" our model and method by splitting the data into training and testing.

The “improvement” in the model is to also include MCAT and Sex as explanatory variables.

```
glm.fits.m=glm(Acceptance~MCAT+GPA+Sex,data=data,
                family=binomial)
summary(glm.fits.m)

##
## Call:
## glm(formula = Acceptance ~ MCAT + GPA + Sex, family = binomial,
##      data = data)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -23.9851    6.9685  -3.442 0.000578 ***
## MCAT         0.1809    0.1080   1.675 0.093946 .
## GPA          5.1392    1.8508   2.777 0.005491 **
## SexM        -1.2580    0.7303  -1.723 0.084965 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 75.791 on 54 degrees of freedom
## Residual deviance: 50.786 on 51 degrees of freedom
## AIC: 58.786
##
## Number of Fisher Scoring iterations: 5
cor(MCAT, GPA)

## [1] 0.5414202
```

The GPA is still the most significant. Let's create the training subset and fit the model on it.

```
n=length(GPA)
set.seed(1)
train=sample(n,floor(n/2))
glm.fits.tr=glm(Acceptance~MCAT+GPA+Sex,data=data,
                 family=binomial, subset=train)
summary(glm.fits.tr)

##
## Call:
## glm(formula = Acceptance ~ MCAT + GPA + Sex, family = binomial,
##      data = data, subset = train)
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -22.0670   10.1235  -2.180  0.0293 *
## MCAT         0.2841    0.1563   1.818  0.0691 .
## GPA          3.3615    2.3373   1.438  0.1504
## SexM        -0.3355    0.9636  -0.348  0.7277
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
```

```

##      Null deviance: 37.096  on 26  degrees of freedom
## Residual deviance: 26.760  on 23  degrees of freedom
## AIC: 34.76
##
## Number of Fisher Scoring iterations: 5

```

How did we do on the training set?

```

glm.probs.tr=predict(glm.fits.tr, type="response")
glm.pred.tr=rep("yes",floor(n/2))
glm.pred.tr[glm.probs.tr<.5]="no"
#Confusion Matrix Below
tab.tr=table(glm.pred.tr,Acceptance[train])
tab.tr

##
## glm.pred.tr  0   1
##           no   9   2
##           yes  3 13
good.tr=(tab.tr[1,1]+tab.tr[2,2])/sum(tab.tr)
good.tr

## [1] 0.8148148

```

Next, we predict using the above model on the validation set, i.e., the complement of `train`.

```

val=data[-train,]
glm.probs.v=predict(glm.fits.tr, data=val ,type="response")
glm.probs.v

##        4       39       1       34       23       43       14
## 0.84563130 0.64594069 0.71038327 0.42415987 0.91597345 0.81236250 0.23569238
##        18      33      21      46      42      10       7
## 0.64035661 0.72163122 0.88615598 0.64369677 0.26001509 0.66116120 0.88468156
##         9      15      45      37      25      52      38
## 0.58601074 0.09810083 0.25359969 0.43158911 0.22701696 0.31511157 0.04892427
##        40      47      20       3       6      50
## 0.25888331 0.99271563 0.98674824 0.13773764 0.58168610 0.79403402

```

How did we do?

```

glm.pred.v=rep("yes", n=floor(n/2))
glm.pred.v[glm.probs.v<.5]="no"
#Confusion Matrix Below
tab.v=table(glm.pred.v,val$Acceptance)
tab.v

##
## glm.pred.v  0   1
##           no   4   7
##           yes  9   8
good.pred=(tab.v[1,1]+tab.v[2,2])/sum(tab.v)
good.pred

## [1] 0.4285714

```