

PCA and K-Means Clustering

Milica Cudina

We consider the `seeds` data set. This data set contains measurements of seeds. First, we import the data set.

```
seeds<-read.csv("seeds_dataset.csv")
seeds

##      X15.26 X14.84 X0.871 X5.763 X3.312 X2.221 X5.22
## 1    14.88  14.57 0.8811  5.554  3.333  1.018  4.956
## 2    14.29  14.09 0.9050  5.291  3.337  2.699  4.825
## 3    13.84  13.94 0.8955  5.324  3.379  2.259  4.805
## 4    16.14  14.99 0.9034  5.658  3.562  1.355  5.175
## 5    14.38  14.21 0.8951  5.386  3.312  2.462  4.956
## 6    14.69  14.49 0.8799  5.563  3.259  3.586  5.219
## 7    14.11  14.10 0.8911  5.420  3.302  2.700    NA
## 8    16.63  15.46 0.8747  6.053  3.465  2.040  5.877
## 9    16.44  15.25 0.8880  5.884  3.505  1.969  5.533
## 10   15.26  14.85 0.8696  5.714  3.242  4.543  5.314
## 11   14.03  14.16 0.8796  5.438  3.201  1.717  5.001
## 12   13.89  14.02 0.8880  5.439  3.199  3.986  4.738
## 13   13.78  14.06 0.8759  5.479  3.156  3.136  4.872
## 14   13.74  14.05 0.8744  5.482  3.114  2.932  4.825
## [ reached 'max' / getOption("max.print") -- omitted 195 rows ]

dim(seeds)

## [1] 209    7
```

We immediately see that there are missing data points. I will choose to omit those rows in the future analysis.

```
seeds=na.omit(seeds)
dim(seeds)
```

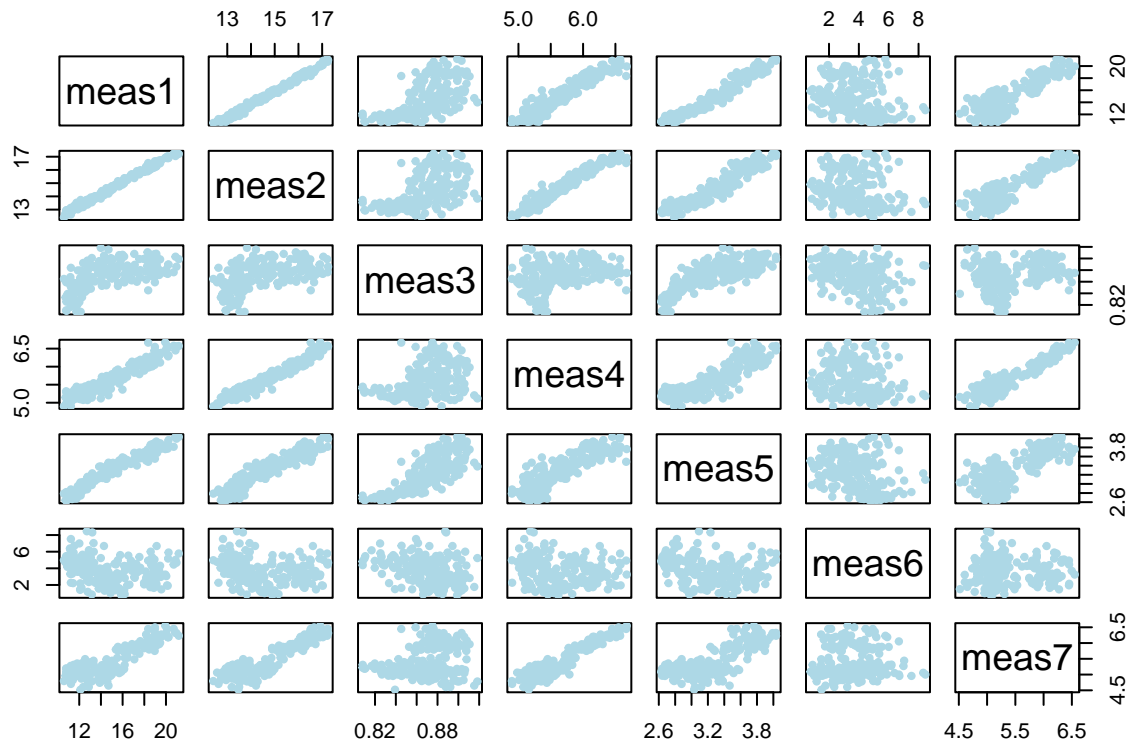
```
## [1] 202    7
```

The columns are not labeled, so I will add names for aesthetic reasons.

```
colnames(seeds)=c("meas1", "meas2", "meas3", "meas4", "meas5", "meas6", "meas7")
attach(seeds)
```

Here is the usual EDA.

```
plot(seeds, col="lightblue", pch=20)
```



We do recognize the strong relationship between some measurements. We could perform K-means clustering on the entire data set.

```
km.out <- kmeans(seeds, 2, nstart = 20)
km.out

## K-means clustering with 2 clusters of sizes 127, 75
##
## Cluster means:
##      meas1      meas2      meas3      meas4      meas5      meas6      meas7
## 1 12.94039 13.70441 0.8629654 5.352441 3.023173 3.893568 5.099126
## 2 18.23107 16.08280 0.8844160 6.126080 3.671987 3.420161 5.965347
##
## Clustering vector:
##  1  2  3  4  5  6  8  9 10 11 12 13 14 15 16 17 18 19 20 21
##  1  1  1  2  1  1  2  2  1  1  1  1  1  1  1  2  1  1  1  1
## 22 23 24 25 26 27 28 29 30 31 32 33 34 36 37 38 39 40 41 42
##  2  1  1  2  1  1  1  1  1  1  1  1  1  2  2  1  1  1  1  1
## 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 61 62 63
##  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1
## 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83
##  1  1  1  1  1  1  2  2  2  2  2  2  2  2  2  2  2  2  2  2
## 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102 103
##  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2  2
## [ reached getOption("max.print") -- omitted 102 entries ]
##
## Within cluster sum of squares by cluster:
## [1] 641.9963 318.6342
## (between_SS / total_SS =  63.6 %)
##
## Available components:
```

```
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
pre.clustering=km.out$cluster
```

However, visualization is challenging in a 7–dimensional space.

So, we want to start with PCA for this data set to reduce the dimension. Let's import the requisite library.

```
library(stats)
```

Let's look at the principal components analysis.

```
pr.out=prcomp(seeds,scale=TRUE)
pr.out$center
```

```
##      meas1      meas2      meas3      meas4      meas5      meas6      meas7
## 14.9047525 14.5874752  0.8709297  5.6396832  3.2640693  3.7177980  5.4207426
```

```
pr.out$scale
```


```
##      meas1      meas2      meas3      meas4      meas5      meas6      meas7
##  2.9270018  1.3129940  0.0233935  0.4445863  0.3793446  1.5045096  0.4937631
```

```
pr.out$rotation
```

```
##           PC1           PC2           PC3           PC4           PC5           PC6
## meas1  0.4438209 -0.02837450  0.02522182 -0.1997502  0.1959546  0.42707585
## meas2  0.4409472 -0.08383445 -0.06056566 -0.3028393  0.1611726  0.47786506
## meas3  0.2785985  0.51914024  0.63899850  0.3405375 -0.3227615  0.13865938
## meas4  0.4238398 -0.19999038 -0.21246198 -0.2464757 -0.7690757 -0.28425740
## meas5  0.4324202  0.11208856  0.21551095 -0.2064352  0.4714107 -0.69824260
## meas6 -0.1193594 -0.72725351  0.66831944 -0.0918059 -0.0383849  0.01704236
## meas7  0.3874240 -0.37576910 -0.22167690  0.8003926  0.1275587 -0.03727380
##
##           PC7
## meas1  0.735451989
## meas2 -0.669944665
## meas3 -0.072041702
## meas4  0.046835025
## meas5 -0.040220518
## meas6 -0.003552126
## meas7 -0.035646645
```

What does the biplot tell us?

```
biplot(pr.out,scale=0, cex=0.5, xlab=rep("*", length(meas1)),
       col=c("lightblue", "red"))
```




Let's look at the variance explained.

```
pr.var=pr.out$sdev^2
pve=pr.var/sum(pr.var)
pve
```

```
## [1] 0.7207429722 0.1665285024 0.0994654055 0.00972229829 0.0026723103
## [6] 0.0007529615 0.0001148653
```

```
par(mfrow = c(1, 2))
plot(pve,xlab="Principal Component",
     ylab="Proportion of Variance Explained", col="peru",
     pch=20, ylim=c(0,1),type='b')
plot(cumsum(pve),xlab="Principal Component",
     ylab="Cumulative Proportion of Variance Explained",
     col="peru", pch=20, ylim=c(0,1),type='b')
```



The values corresponding to our rows in terms of principal components are available through `pr.out$x`.

`pr.out$x`

##	PC1	PC2	PC3	PC4	PC5
## 1	-0.04207784	1.94468075	-0.632127428	-0.424848174	0.0385772091
## 2	-0.49042490	1.91804621	0.971251764	-0.097261002	-0.0063964294
## 3	-0.62362012	1.94663985	0.536415581	-0.216966866	0.0773032846
## 4	1.06088821	2.09183935	0.100261549	-0.130928172	0.0194619872
## 5	-0.37072106	1.65455973	0.472378239	-0.087443426	-0.0045658223
## 6	-0.18519061	0.45760937	0.313464401	-0.105994075	-0.0725094802
## 8	1.71369945	0.34847250	-0.955896688	0.239391431	-0.1340264789
## 9	1.39295855	1.04282786	-0.358166002	0.013032605	-0.1010270808
## 10	0.02266578	-0.40731929	0.321063526	-0.356726069	-0.1302589989
## 11	-0.60772047	1.58685848	-0.390767934	-0.127676023	-0.0175404493
## 12	-0.96362326	0.88588011	0.968306548	-0.527798816	-0.2900418275
## 13	-0.94927834	0.89407142	0.153716027	-0.435350237	-0.1918784475
## 14	-1.04227573	0.98242904	0.018047184	-0.494695701	-0.2394084500
## 15	-0.54878468	1.06352121	1.458281375	-0.437619311	-0.0423886542
##	PC6	PC7			
## 1	-0.01728263413	-0.00503191898			
## 2	0.05331998378	-0.00460471997			
## 3	-0.22512186236	-0.01037567716			
## 4	-0.04911991102	-0.00134456779			
## 5	0.02414801380	-0.00897167289			
## 6	0.05845540204	-0.02451329546			
## 8	-0.09592381419	-0.03005666416			
## 9	-0.06165135197	-0.00864637487			
## 10	0.15000942325	-0.02466840104			
## 11	0.02223897884	-0.00791191004			
## 12	0.04924536861	0.01642390353			

```
## 13  0.00986276075  0.00674361030
## 14  0.06812257211  0.01505867892
## 15  0.12161224997 -0.00220489260
## [ reached getOption("max.print") -- omitted 188 rows ]
```

We can isolate the first two principal components as follows:

```
pc1=pr.out$x[,1]
pc2=pr.out$x[,2]
```

Let's look at the scatterplot of these two vectors.

```
plot(pc1, pc2, col="grey",pch=20)
```

